



Deep Learning

Andrea Asperti

DISI - Department of Informatics: Science and Engineering
University of Bologna
Mura Anteo Zamboni 7, 40127, Bologna, ITALY
andrea.asperti@unibo.it

Q: What is Deep Learning?



Deep Learning in a nutshell



Q: What is Deep Learning?

A: A branch of **Machine Learning**



Deep Learning in a nutshell



Q: What is Deep Learning?

A: A branch of **Machine Learning**

Q: To what kind of problems it can be applied?



Deep Learning in a nutshell



Q: What is Deep Learning?

A: A branch of **Machine Learning**

Q: To what kind of problems it can be applied?

A: To **all problems** suitable for Machine Learning



Deep Learning in a nutshell



Q: What is Deep Learning?

A: A branch of **Machine Learning**

Q: To what kind of problems it can be applied?

A: To **all problems** suitable for Machine Learning

Q: Are there problems where DL particularly excels?





Q: What is Deep Learning?

A: A branch of **Machine Learning**

Q: To what kind of problems it can be applied?

A: To **all problems** suitable for Machine Learning

Q: Are there problems where DL particularly excels?

A: All problems involving **myriads of features**, e.g.
images/speech/text processing, recommendation systems, ...





Q: What is Deep Learning?

A: A branch of **Machine Learning**

Q: To what kind of problems it can be applied?

A: To **all problems** suitable for Machine Learning

Q: Are there problems where DL particularly excels?

A: All problems involving **myriads of features**, e.g.
images/speech/text processing, recommendation systems, ...





Q: What is the basic technique for DL?



Deep Learning FAQ (2)



Q: What is the basic technique for DL?

A: Neural Networks



Deep Learning FAQ (2)



Q: What is the basic technique for DL?

A: Neural Networks

Q: Why “Deep”?



Q: What is the basic technique for DL?

A: Neural Networks

Q: Why “Deep”?

A: - because it exploits Deep Neural Networks, composed by many **nested layers** of neurons





Q: What is the basic technique for DL?

A: Neural Networks

Q: Why “Deep”?

- A:**
- because it exploits Deep Neural Networks, composed by many **nested layers** of neurons
 - because it exploits **deep features** of data, that is features extracted from other features

ML recap

What is Machine Learning about?

There are problems that are difficult to address with traditional programming techniques:

- ▶ classify a document according to some criteria (e.g. spam, sentiment analysis, ...)
- ▶ compute the probability that a credit card transaction is fraudulent
- ▶ recognize an object in some image (possibly from an unusual viewpoint, in new lighting conditions, in a cluttered scene)
- ▶ ...

Typically the result is a weighted combination of a large number of parameters, each one contributing to the solution in a small degree.

The Machine Learning approach

Suppose to have a set of input-output pairs (**training set**)

$$\{\langle x_i, y_i \rangle\}$$

the problem consists in guessing the map $x_i \mapsto y_i$




The M.L. approach:

- describe the problem with a **model** depending on some parameters Θ (i.e. choose a parametric class of functions)
- define a **loss function** to compare the results of the model with the expected (experimental) values
- **optimize** (fit) the parameters Θ to reduce the loss to a minimum

Example: a regression problem

You have some points on the plane and you want to fit a line through them

- Step 1** Fix a parametric class of models.
For instance linear functions $y = ax + b$;
 a and b are the **parameters** of the model
- Step 2** Fix a way to decide when a line is better than another (loss function)
For instance, mean square error (mse)
- Step 3** Try to tune the parameters in order to reduce the loss (training).



Why Learning?

Machine Learning problems are in fact **optimization problems!**
So, why talking about learning?

Why Learning?

Machine Learning problems are in fact **optimization problems!**
So, why talking about learning?

The point is that the solution to the optimization problem is not given in an analytical form (often there is no closed form solution).

Why Learning?

Machine Learning problems are in fact **optimization problems!**
So, why talking about learning?

The point is that the solution to the optimization problem is not given in an analytical form (often there is no closed form solution).

So, we use **iterative** techniques (typically, gradient descent) to progressively approximate the result.

Why Learning?

Machine Learning problems are in fact **optimization problems!**
So, why talking about learning?

The point is that the solution to the optimization problem is not given in an analytical form (often there is no closed form solution).

So, we use **iterative** techniques (typically, gradient descent) to progressively approximate the result.

This form of iteration over data can be understood as a way of progressive learning of the objective function based on the experience of past observations.

Using gradients

Goal: minimize a loss function E over (fixed) training samples:


$$\Theta(w) = \sum_i E(o(w, x_i), y_i)$$

See how it changes according to small perturbations $\Delta(w)$ of the parameters w : this is the **gradient**

$$\nabla_w[\theta] = \left[\frac{\partial \theta}{\partial w_1}, \dots, \frac{\partial \theta}{\partial w_n} \right]$$

of Θ w.r.t. w .

The gradient is a **vector** pointing in the direction of **steepest ascent**.




A bit of taxonomy

Different types of Learning Tasks

- **supervised learning:**

inputs + outputs (labels)


- classification
- regression



- **unsupervised learning:**

just inputs

- clustering
- component analysis
- anomaly detection
- autoencoding



- **reinforcement learning**


actions and rewards

- learning long-term gains
- planning




Classification vs. Regression

Two forms of supervised learning: $\{\langle x_i, y_i \rangle\}$



classification



regression


y is discrete: $y \in \{\bullet, +\}$

y is (conceptually) continuous


Many different techniques

- **Different ways to define the models:**

- decision trees
- linear models
- neural networks
- ...




decision tree




neural net

- **Different error (loss) functions:**

- mean squared errors
- logistic loss
- cross entropy
- cosine distance
- maximum margin
- ...




mean squared errors



maximum margin




Neural Networks




Neural Network

A network of (artificial) neurons



Each neuron takes multiple inputs and produces a single output (that can be passed as input to many other neurons).

The artificial neuron







Each neuron (!) implements a logistic regressor


$$\sigma(wx + b)$$

Different activation functions

The activation function is responsible for threshold triggering.




The cortical neuron



- ▶ the **dendritic tree** of the cell collects inputs from other neurons, that get summed together
- ▶ when a **triggering threshold** is exceeded, the Axon Hillock generate an impulse that get transmitted through the axon to other neurons.

Artificial Neural Networks (ANN)



Slide credit : Andrew L. Nelson



Some figures for human brains

- ▶ number of neurons: $\sim 2 \cdot 10^{10}$
- ▶ switching time for neuron: 1 – 5 ms. (**slow!**)
- ▶ synapses (connections) per neuron: $\sim 10^4$ – 10^5
- ▶ time to process an image: 100 ms.

not too deep (< 100)
very high parallelism



Motivations behind neural computation

- ▶ to understand, via simulation, how the brain works
- ▶ to investigate a different paradigm of computation
very far from traditional programming languages
- ▶ **to solve practical problems difficult to address with algorithmic techniques**
useful even if the brain works in a different way



Network topologies

Feed-forward and recurrent networks


If the network is acyclic, it is called a **feed-forward** network.

If it has cycles it is called **recurrent**.



Layers


In a feed-forward network, neurons are usually organized in **layers**.



If there is more than one hidden layer the network is **deep**, otherwise it is called a **shallow** network.

Main layers in feed-forward networks: dense layer

Dense layer: each neuron at layer $k-1$ is connected to **each** neuron at layer k .



A single neuron:

$$I^n \cdot W^n + B^1 = O^1$$

the operation can be **vectorized** to produce m outputs in parallel:


$$I^n \cdot W^{n \times m} + B^m = O^m$$

- ▶ dense layers usually work on **flat** (unstructured) inputs
- ▶ the order of elements in input is **irrelevant**



Main layers in feed-forward networks: convolutional layer

Convolutional layer: each neuron at layer $k - 1$ is connected via a parametric kernel to a fixed subset of neurons at layer k . The kernel is convolved over the whole previous layer.



1. move the kernel K over a portion M of the input of equal size
2. compute the dot product $M \cdot K$ and possibly add a bias
3. shift the kernel and repeat

The dimension of the output only depends from the number of times the kernel is applied.

Input is **structured**, and the structure is reflected in the output.

Parameters and hyper-parameters

The weights W_k are the **parameters** of the model: they are learned during the training phase.

The number of neurons and the way they are connected together are **hyper-parameters**: they are chosen by the user and fixed before training may start.

Other important hyper-parameters govern training such as **learning rate**, **batch-size**, number of **epochs** and many others.



Features and deep features

Any individual measurable property of data useful for the solution of a specific task is called a **feature**.

Examples:

- ▶ **Medical Diagnosis:** information about the patient (age, clinical history, ...), symptoms, physical examination, results of medical tests, ...
- ▶ **Meteo forecasting:** humidity, pressure, temperature, wind, rain, snow, ...
- ▶ **Image Processing:** raw pixels, combination of adjacent pixels, ...


- **Signals:** raw input collected from sensors
- **Data:** meaningful but not focused
- **Features:** meaningful and focused

Deep learning, in deeper sense

Discovering good features is a **complex task**.

Why not delegating the task to the machine, **learning** them?

Deep learning exploits a *hierarchical organization* of the learning model, allowing complex features to be computed in terms of simpler ones, through non-linear transformations.



Each layer synthesizes new features in terms of the previous ones.




- **Knowledge-based systems:** take an expert, ask him how he solves a problem and try to mimic his approach by means of logical rules

- **Knowledge-based systems:** take an expert, ask him how he solves a problem and try to mimic his approach by means of logical rules
- **Traditional Machine-Learning:** take an expert, ask him what are the features of data relevant to solve a given problem, and let the machine learn the mapping

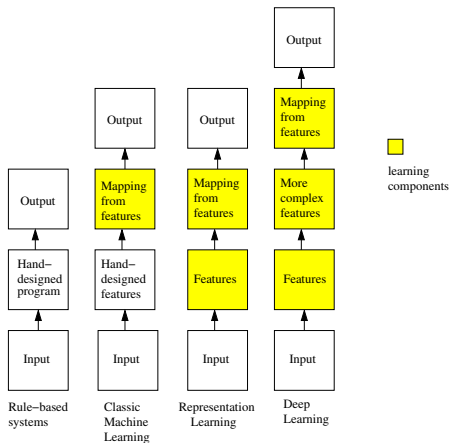
- **Knowledge-based systems:** take an expert, ask him how he solves a problem and try to mimic his approach by means of logical rules
- **Traditional Machine-Learning:** take an expert, ask him what are the features of data relevant to solve a given problem, and let the machine learn the mapping
- **Deep-Learning:** get rid of the expert

Relations between research areas



Picture from “Deep Learning” by Y.Bengio, I.Goodfellow e A.Courville, MIT Press.

Components trained to learn



Picture from “Deep Learning” by Y.Bengio, I.Goodfellow e A.Courville, MIT Press.

Diving into DL



A first example

[demo]



Understanding DL

- understand the **different layers**, and their purpose
- understand how layers can be organized in **relevant architectures**
- understand the different possible **applications** of DL, and their specific solutions
- understand the main **issues, problems** and **costs**

- TensorFlow/Keras, Google Brain
- PyTorch, Facebook
- MXNET, Apache

We shall mostly use Keras.

Historical remarks - Legacy

Legacy

1958	perceptron
1975	backpropagation
1980	convolutional layers
1992	Max-pooling
1997	LSTM
...	...

Extremely slow progress

Neural Networks played a marginal role in AI



The Deep Learning revolution

2011	Google Brain foundation	2017	Mask-RCNN
2012	ReLU and Dropout	2017	PPO
2012	ImageNet Competition	2018	Transformers
2013	DQN	2018	BERT/GPT
2014	GANs	2018	Soft Actor Critic
2014	Attention	2019	AlphaStar
2014	Inception v1	2020	Vision Transformers
2015	Tensorflow release	2020	OpenAI Jukebox
2015	Keras release	2021	MXNet release
2015	AlphaGo	2021	TFP release
2015	Batchnormalization	2022	Keras-CV
2015	YOLO v1	2022	Dall-E 2, Imagen
2015	OpenAI foundation	2022	ChatGPT
2016	Residual connections	2023	LLaMA-2
2017	PyTorch release	2023	ChatGPT-4

Just to mention a few milestones ...

The Deep Learning revolution

Frameworks and Libraries

2011	Google Brain foundation	2017	Mask-RCNN
2012	ReLU and Dropout	2017	PPO
2012	ImageNet Competition	2018	Transformers
2013	DQN	2018	BERT/GPT
2014	GANs	2018	Soft Actor Critic
2014	Attention	2019	AlphaStar
2014	Inception v1	2020	Vision Transformers
2015	Tensorflow release	2020	OpenAI Jukebox
2015	Keras release	2021	MXNet release
2015	AlphaGo	2021	TFP release
2015	Batchnormalization	2022	Keras-CV
2015	YOLO v1	2022	Dall-E 2, Imagen
2015	OpenAI foundation	2022	ChatGPT
2016	Residual connections	2023	LLaMA-2
2017	PyTorch release	2023	ChatGPT-4

The Deep Learning revolution

Some popular technical improvements

2011	Google Brain foundation	2017	Mask-RCNN
2012	ReLU and Dropout	2017	PPO
2012	ImageNet Competition	2018	Transformers
2013	DQN	2018	BERT/GPT
2014	GANs	2018	Soft Actor Critic
2014	Attention	2019	AlphaStar
2014	Inception v1	2020	Vision Transformers
2015	Tensorflow release	2020	OpenAI Jukebox
2015	Keras release	2021	MXNet release
2015	AlphaGo	2021	TFP release
2015	BatchNormalization	2022	Keras-CV
2015	YOLO v1	2022	Dall-E 2, Imagen
2015	OpenAI foundation	2022	ChatGPT
2016	Residuality & Resnet	2023	LLaMA-2
2017	PyTorch release	2023	ChatGPT-4

The Deep Learning revolution

Important Architectures

2011	Google Brain foundation	2017	Mask-RCNN
2012	ReLU and Dropout	2017	PPO
2012	ImageNet Competition	2018	Transformers
2013	DQN	2018	BERT/GPT
2014	GANs	2018	Soft Actor Critic
2014	Attention	2019	AlphaStar
2014	Inception v1	2020	Vision Transformes
2015	Tensorflow release	2020	OpenAI Jukebox
2015	Keras release	2021	MXNet release
2015	AlphaGo	2021	TFP release
2015	BatchNormalization	2022	Keras-CV
2015	YOLO v1	2022	Dall-E 2, Imagen
2015	OpenAI foundation	2022	ChatGPT
2016	Residuality & Resnet	2023	LLaMA-2
2017	PyTorch release	2023	ChatGPT-4

The Deep Learning revolution

Deep Reinforcement Learning


2011	Google Brain foundation	2017	Mask-RCNN
2012	ReLU and Dropout	2017	PPO
2012	ImageNet Competition	2018	Transformers
2013	DQN	2018	BERT/GPT
2014	GANs	2018	Soft Actor Critic
2014	Attention	2019	AlphaStar
2014	Inception v1	2020	Vision Transformes
2015	Tensorflow release	2020	OpenAI Jukebox
2015	Keras release	2021	MXNet release
2015	AlphaGo	2021	TFP release
2015	BatchNormalization	2022	Keras-CV
2015	YOLO v1	2022	Dall-E 2, Imagen
2015	OpenAI foundation	2022	ChatGPT
2016	Residuality & Resnet	2023	LLaMA-2
2017	PyTorch release	2023	ChatGPT-4




Generative AI and Large Language Models

2011	Google Brain foundation	2017	Mask-RCNN
2012	ReLU and Dropout	2017	PPO
2012	ImageNet Competition	2018	Transformers
2013	DQN	2018	BERT/GPT
2014	GANs	2018	Soft Actor Critic
2014	Attention	2019	AlphaStar
2014	Inception v1	2020	Vision Transformers
2015	Tensorflow release	2020	OpenAI Jukebox
2015	Keras release	2021	MXNet release
2015	BatchNormalization	2021	TFP release
2015	YOLO v1	2022	Keras-CV
2015	OpenAI foundation	2022	Dall-E 2, Imagen
2016	Residuality & Resnet	2022	ChatGPT
2017	PyTorch release	2023	LLaMA-2
		2023	ChatGPT-4

The situation at the beginning of the century



The deep learnig era



See my [blog](#) for a short historical perspective.



- ▶ structure of the course
- ▶ books, tutorials and blogs
- ▶ software
- ▶ examination
- ▶ office hours

Frontal lessons intermixed with demos + labs

- Domains of application of Deep Learning
- Expressiveness
- Backpropagation
- Convolutional Networks
- Understanding CNNs
- Object Detection and Segmentation
- Autoencoders
- Generative Adversarial Networks
- Recurrent Networks
- LSTM, Attention, Transformers
- Reinforcement Learning

- ▶ Dive into deep learning (D2L)
- ▶ Y.Bengio, I.Goodfellow and A.Courville. **Deep Learning**, MIT Press to appear.

Possible to study on online material (fast updating):

- ▶ [Tensorflow tutorials](#)
- ▶ [Towards data science](#)
- ▶ [Keras blog](#). By F.Chollet.
- ▶ [Machine learning tutorial with Python](#)
- ▶ [Deep Learning Tutorial](#). LISA lab. University of Montreal.
- ▶ a lot of interesting lessons and seminars on youtube
- ▶ a lot of material on github
- ▶ ...

The State of the Art site! (papers with code)

- Tensor flow dataset
- Kaggle Datasets
- Many standard datasets for image processing: Pascal VOC, Coco, ...
- Face detection and recognition: CelebA, Labeled Faces in the Wild, ...
- Biomedical challenges
- Amazon Datasets
- ...

At each exam session you will receive a project assignement that you are supposed to complete in **7 days**.

You are supposed to deliver:

- ▶ the code source (Keras/TensorFlow) in the form of a **single, commented** pyhton notebook

The work will be evaluated according to

1. 80% : comparative evaluation of results (measured in an objective way according to given metrics);
2. 20% : descriptive quality of the notebook

You may possibly integrate the grade with an oral examination.



Office hours: On appointment

Prof. Andrea Asperti

andrea.asperti@unibo.it

Via Zanolini 41 - Ex Veneta Station

Tutors:

- Salvatore Fiorilla: salvatore.fiorilla@unibo.it
- Fabio Merizzi: fabio.merizzi@unibo.it

